


Putting it all together – Final Review

CS 475: Concurrent & Distributed Systems (Fall 2021)

Yue Cheng

Back in Lec-1...

Distributed systems: What?



- Multiple cooperating computers
 - Connected by a network
 - Doing something together
- Lots of critical infrastructure are distributed

Distributed systems: Why?

- Or, why not 1 computer to rule them all?
- Failure
- Limited computation/storage
- Physical location

Distributed systems: Why?

- Or, why not 1 computer to rule them all?

- Failure ➤ Fault tolerance

- Limited computation/storage ➤ Scalability

- Physical location ➤ Availability, low latency

Goals of “distributed systems”

- Service with higher-level abstractions/interface
 - E.g., key-value store, programming model, ...
- High complexity
 - Scalable (scale-out)
 - Reliable (fault-tolerant)
 - Well-defined semantics (consistent)
- Do “heavy lifting” so app developers don’t need to

Theme

- Fundamental building blocks
- Abstractions and programming models
- Production system designs

Theme

- Fundamental building blocks
- Abstractions and programming models
- Production system designs

Fundamental building blocks

- Remote procedure calls (RPCs)

Fundamental building blocks

- Remote procedure calls (RPCs)
- Time & clocks

Fundamental building blocks

- Remote procedure calls (RPCs)
- Time & clocks
- Consensus algorithms

Fundamental building blocks

- Remote procedure calls (RPCs)
- Time & clocks
- Consensus algorithms
- Replication, sharding, transactions

Fundamental building blocks

- Remote procedure calls (RPCs)
- Time & clocks
 - **Vector clocks**
- Consensus algorithms
 - **Raft**
- Replication, sharding, transactions
 - **Serializability**

Theme

- Fundamental building blocks
- Abstractions and programming models
- Production system designs

Programming models

- MapReduce
- Spark

Programming models

- MapReduce

- **Spark**

Resilient Distributed Datasets & Spark

- Transformations and actions
- `persist()`
 - Not an action nor a transformation – tell which RDDs should materialize
- PageRank example
 - How iterative PR algorithm works
 - Where to place `persist()` in iterative PR

Theme

- Fundamental building blocks
- Abstractions and programming models
- **Production system designs**

Production system designs

- Amazon Dynamo
- Facebook memcache

Production system designs

- Amazon Dynamo
- **Facebook memcache**

Facebook memcache

- Memcache as a demand-filled, look-aside cache
 - `Read()` and `write()`
- Interesting problems solved in FB's production-scale memcache deployments
 1. Stale set: **a single region** vs. **geographically distributed**
 2. Thundering herds
 3. Incast congestion
 4. Incorporating McSqueal for what?

Final exam

- Thursday, Dec 09th, 7:30 – 10:00am
 - 150 minutes
 - Open-book, open-notes (you may use class notes, papers, and lab materials; you may read them on your laptop, **but you are not allowed to use any network**)
 - Let me know if you need testing center accommodation ASAP (**no guarantee if you send me the form one day before the final exam**)
- Covering topics from lec-1 to lec-17
 - High-level design questions
 - **30%** before midterm **70%** after midterm

Topics

1. Vector clocks
2. Raft
3. Transactions
4. Spark
5. Facebook memcache

**Don't forget to fill out the course
evaluation form**

Good luck! 😊