

Concurrency Control, Recovery, and Locking

CS 475: Concurrent & Distributed Systems (Fall 2021)

Lecture 14

Yue Cheng

Some material taken/derived from:

- Princeton COS-418 materials created by Michael Freedman and Kyle Jamieson.
- MIT 6.824 by Robert Morris, Frans Kaashoek, and Nickolai Zeldovich.

Licensed for use under a Creative Commons Attribution-NonCommercial-ShareAlike 3.0 Unported License.

The transaction

- Definition: A **unit** of work:
 - May consist of **multiple** data accesses or updates
 - Must **commit** or **abort** as a **single atomic unit**
- Transactions can either **commit**, or **abort**
 - When **commit**, all updates performed on database are made permanent, visible to other transactions
 - When **abort**, database restored to a state such that the aborting transaction never executed

Defining properties of transactions

- **Atomicity**: Either **all** constituent operations of the transaction complete successfully, or **none** do
- **Consistency**: Each transaction in isolation preserves a set of **integrity constraints** on the data
- **Isolation**: Transactions' behavior not impacted by presence of **other concurrent transactions**
- **Durability**: The transaction's **effects survive failure** of volatile (memory) or non-volatile (disk) storage

Challenges

1. High transaction **speed requirements**
 - If always `fsync()` to disk for each result on transaction, yields terrible performance
2. **Atomic and durable** writes to disk are difficult
 - In a manner to handle arbitrary crashes
 - Hard disks and solid-state storage use **write buffers** in volatile memory

Today

Techniques for achieving ACID properties

- Write-ahead logging and checkpointing
- Serializability and two-phase locking

What does the system need to do?

- Transaction's properties: **ACID**
 - Atomicity, Consistency, Isolation, Durability
- **Application logic** checks **consistency (C)**
- This leaves **two main goals** for the **system**:
 1. Handle **failures (A, D)**
 2. Handle **concurrency (I)**

Failure model: crash failures

- Standard “crash failure” model:
- Machines are prone to crashes:
 - Disk contents (*non-volatile storage*) **okay**
 - Memory contents (*volatile storage*) **lost**
- Machines don’t misbehave (“Byzantine”)

Account transfer transaction

- Transfers \$10 from account *A* to account *B*

```
transaction transfer(A, B):  
  begin_tx  
  a ← read(A)  
  if a < 10 then abort_tx  
  else write(A, a-10)  
    b ← read(B)  
    write(B, b+10)  
  commit_tx
```


Problem

- Suppose \$100 in A, \$100 in B


```
transaction transfer(A, B):  
begin_tx  
a ← read(A)  
if a < 10 then abort_tx  
else   write(A, a-10)  
       b ← read(B)  
       write(B, b+10)  
       commit_tx
```

- *commit_tx* starts the commit protocol:
 - write(A, \$90) to disk
 - write(B, \$110) to disk
- What happens if **system crash** after first write, but before second write?
 - After recovery: Partial writes, **money is lost**

Lack atomicity in the presence of failures

System architecture

- Smallest unit of storage that can be atomically written to non-volatile storage is called a **page**
- **Buffer manager** moves pages between **buffer pool** (in volatile memory) and disk (in non-volatile storage)



Two design choices

1. **Force** all a transaction's writes to disk **before** transaction commits?

- Yes: *force* policy
- No: *no-force* policy

2. May **uncommitted** transactions' writes **overwrite** committed values on disk?

- Yes: *steal* policy
- No: *no-steal* policy

Performance implications

1. **Force** all a transaction's writes to disk **before** transaction commits?

- Yes: **force** policy

Then **slower disk writes** appear **on the critical path** of a committing transaction

2. May **uncommitted** transactions' writes **overwrite** committed values on disk?

- No: **no-steal** policy

Then buffer manager **loses write scheduling flexibility**

Undo & redo

1. Force all a transaction's writes to disk **before** transaction commits?

- Choose **no: no-force** policy
 - 👉 Need support for **redo**: complete a committed transaction's writes on disk

2. May **uncommitted** transactions' writes **overwrite** committed values on disk?


- Choose **yes: steal** policy
 - 👉 Need support for **undo**: removing the effects of an uncommitted transaction on disk

How to implement undo & redo?

- **Log:** A sequential file that stores information about transactions and system state
 - Resides in **separate, non-volatile storage**
- One entry in the log for each update, commit, abort operation: called a **log record**
- Log record contains:
 - Monotonic-increasing **log sequence number** (LSN)
 - Old value (**before image**) of the item for undo
 - New value (**after image**) of the item for redo

System architecture

- **Buffer pool** (volatile memory) and disk (non-volatile)
- The **log** resides on a **separate** partition or disk (in non-volatile storage)



Write-ahead Logging (WAL)

- Ensures atomicity in the event of system crashes under **no-force/steal** buffer management
1. **Force all log records** pertaining to an updated page into the (non-volatile) log **before any (over)-writes** to page itself
 2. A transaction is not considered committed until **all its log records** (including commit record) are **forced into the log**

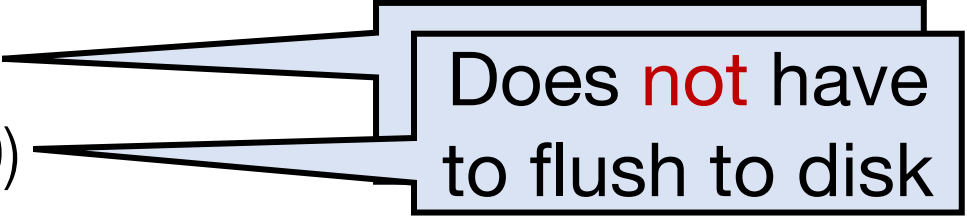
WAL example

force_log_entry(A, old=\$100, new=\$90)

force_log_entry(B, old=\$100, new=\$110)

write(A, \$90)

write(B, \$110)



Does **not** have
to flush to disk

force_log_entry(commit)

- What if the commit log record size > the page size?
- How to ensure **each log record** is written atomically?
 - Write a **checksum** of entire log entry

Goal #2: Concurrency control

Transaction isolation

Two concurrent transactions

```
transaction sum(A, B):  
begin_tx  
a ← read(A)  
b ← read(B)  
print a + b  
commit_tx
```


```
transaction transfer(A, B):  
begin_tx  
a ← read(A)  
if a < 10 then abort_tx  
else write(A, a-10)  
      b ← read(B)  
      write(B, b+10)  
commit_tx
```

Isolation between transactions


- **Isolation:** sum appears to happen either completely before or completely after **transfer**
 - Sometimes called *before-after atomicity*
- *Schedule* for transactions is an ordering of the operations performed by those transactions

Problem for concurrent execution: Inconsistent retrieval

- **Serial execution** of transactions — transfer then sum:



- Concurrent execution resulting in *inconsistent retrieval*, result differing from any serial execution:



Time →
© = commit

Isolation between transactions

- **Isolation:** sum appears to happen either completely before or completely after **transfer**
 - Sometimes called *before-after atomicity*
- Given a schedule of operations:
 - *Is that schedule in some way “equivalent” to a serial execution of transactions?*

Equivalence of schedules

- Two operations from different transactions are *conflicting* if:
 1. They read and write to the same data item
 2. The write and write to the same data item
- Two schedules are *equivalent* if:
 1. They contain the same transactions and operations
 2. They order all conflicting operations of non-aborting transactions in the same way

Conflict serializability


- Ideal isolation semantics: *conflict serializability*
- A schedule is ***conflict serializable*** if it is equivalent to some serial schedule
 - *i.e.*, non-conflicting operations can be **reordered** to get a **serial** schedule

A serializable schedule

- Ideal isolation semantics: *conflict serializability*
- A schedule is ***conflict serializable*** if it is equivalent to some serial schedule
 - *i.e.*, non-conflicting operations can be reordered to get a serial schedule

transfer:


sum:



Time →
© = commit

A **non**-serializable schedule

- Ideal isolation semantics: *conflict serializability*
- A schedule is ***conflict serializable*** if it is equivalent to some serial schedule
 - *i.e.*, **non-conflicting** operations can be **reordered** to get a **serial** schedule



But in a serial schedule, sum's reads
either both before w_A or both after w_B

~~connecting ops~~

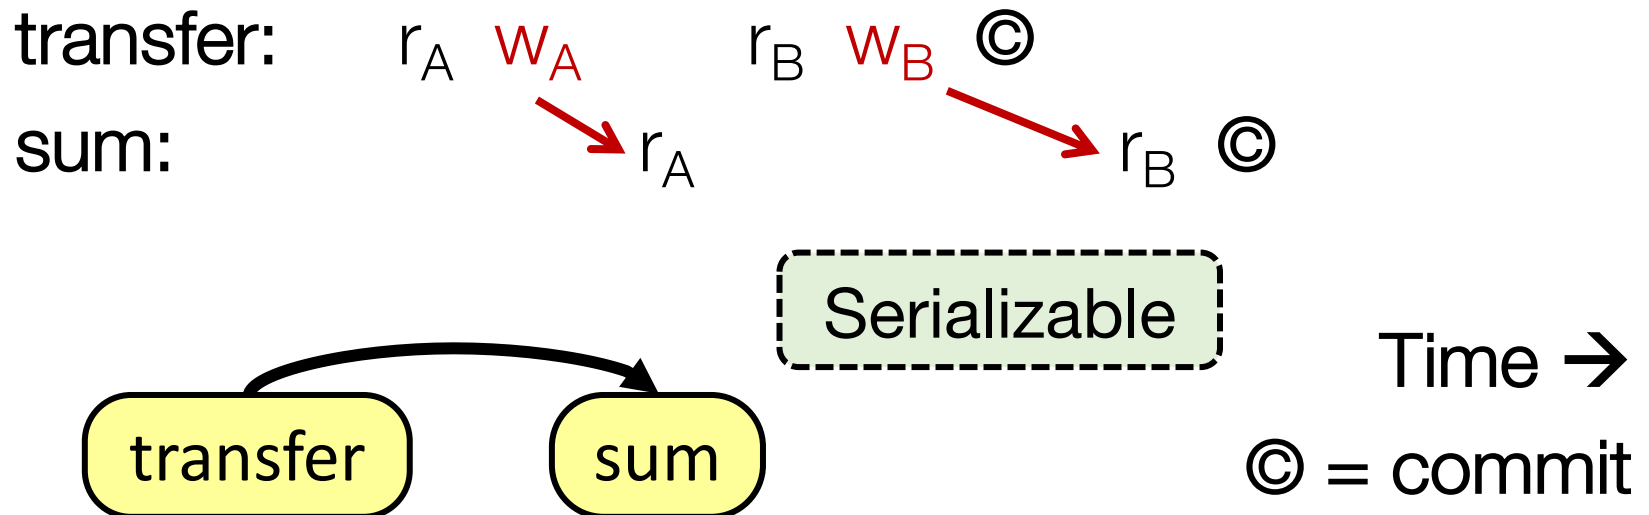
Time →
© = commit

Testing for serializability

- Each node t in the **precedence graph** represents a transaction t
 - Edge from s to t if some action of s **precedes** **and conflicts with** some action of t


Serializable schedule, acyclic graph

- Each node t in the **precedence graph** represents a transaction t
 - Edge from s to t if some action of s **precedes** and **conflicts with** some action of t




Non-serializable schedule, cyclic graph

- Each node t in the **precedence graph** represents a transaction t
 - Edge from s to t if some action of s **precedes** and **conflicts with** some action of t



Non-serializable



Time →
© = commit

Testing for serializability

- Each node t in the **precedence graph** represents a transaction t
 - Edge from s to t if some action of s **precedes** **and conflicts with** some action of t

In general, a schedule is conflict-serializable if and only if its **precedence graph** is **acyclic**

How to ensure a serializable schedule?

- Locking-based approaches
- **Strawman 1: Big Global Lock**
 - Acquire the lock when transaction starts
 - Release the lock when transaction ends

Results in a serial transaction schedule
at the **cost of performance**

Locking

- Locks maintained by **transaction manager**
 - Transaction requests lock **for a data item**
 - Transaction manager **grants** or **denies** lock
- Lock types
 - **Shared**: Need to have before read object
 - **Exclusive**: Need to have before write object

	Shared (S)	Exclusive (X)
Shared (S)	Yes	No
Exclusive (X)	No	No

How to ensure a serializable schedule?

- **Strawman 2:** Grab locks independently, for each data item (e.g., bank accounts A and B)



Permits this non-serializable interleaving

Time →

© = commit

$\blacktriangle / \triangle = \text{eXclusive- / Shared-lock}$; $\blacktriangle / \triangle = \text{X- / S-unlock}$

Two-phase locking (2PL)

- **2PL rule:** Once a transaction has **released** a lock it is **not allowed to obtain** any other locks
- A **growing phase** when transaction acquires locks
- A **shrinking phase** when transaction releases locks
- In practice:
 - Growing phase is the entire transaction
 - Shrinking phase is during commit

2PL allows only serializable schedules

- **2PL rule:** Once a transaction has **released** a lock it is **not allowed to obtain** any other locks

transfer: $\blacktriangle_A r_A w_A \blacktriangle_A$

sum: $\triangle_A r_A \triangle_A \triangle_B r_B \triangle_B \textcircled{\text{C}}$

2PL precludes this non-serializable interleaving

Time →

© = commit

▲ / △ = X- / S-lock; ▼ / ▽ = X- / S-unlock

2PL and transaction concurrency

- **2PL rule:** Once a transaction has **released** a lock it is **not allowed to obtain** any other locks

transfer: $\triangleleft_A r_A$ $\blacktriangleleft_A w_A$ $\triangleleft_B r_B$ $\blacktriangleleft_B w_B * \textcircled{C}$
 sum: $\triangleleft_A r_A$ $\triangleleft_B r_B * \textcircled{C}$

2PL **permits** this **serializable, interleaved** schedule

Time \rightarrow

\textcircled{C} = commit

$\blacktriangleleft / \triangleleft$ = X- / S-lock; $\blacktriangleright / \triangleright$ = X- / S-unlock; $*$ = release all locks

2PL doesn't exploit all opportunities for concurrency

- **2PL rule:** Once a transaction has **released** a lock it is **not allowed to obtain** any other locks

transfer: r_A w_A r_B w_B ©
sum: r_A r_B ©

2PL **precludes** this **serializable, interleaved** schedule

Time →

© = commit

(locking not shown)

Issues with 2PL

- What if a lock is unavailable? Is **deadlock** possible?
 - Yes; but a central controller can detect deadlock cycles and **abort involved transactions**
- The **phantom problem**
 - Database has fancier ops than key-value store
 - T1: begin_tx; update employee (set salary = $1.1 \times \text{salary}$) where dept = “CS”; commit_tx
 - T2: insert into employee (“carol”, “CS”)
 - Even if they lock individual data items, could result in **non-serializable execution**

Serializability vs. linearizability

- **Linearizability**: a guarantee about **single** operations on **single** objects
 - Once write completes, all later reads (by wall clock) should reflect that write
- **Serializability** is a guarantee about **transactions** over **one or more** objects
 - Doesn't impose real-time constraints
- Linearizability + serializability = ***strict serializability***
 - Transaction behavior equivalent to some serial execution
 - And that serial execution agrees with real-time

Today

Techniques for achieving ACID properties

- Write-ahead logging and check-pointing → A, D
- Serializability and two-phase locking → I