

# 2PL and OCC

*CS 475: Concurrent & Distributed Systems (Fall 2021)*

Lecture 15

Yue Cheng

Some material taken/derived from:

- Princeton COS-418 materials created by Michael Freedman and Kyle Jamieson.
- MIT 6.824 by Robert Morris, Frans Kaashoek, and Nickolai Zeldovich.

Licensed for use under a Creative Commons Attribution-NonCommercial-ShareAlike 3.0 Unported License.

# Recap: Transaction serializability

## Serializability:

Execution of a set of transactions over multiple items is equivalent to **some serial execution** of transactions

# Some new terms

**Lost update:** the result of a transaction is overwritten by another transaction

# Some new terms

**Lost update:** the result of a transaction is overwritten by another transaction

**Dirty read:** uncommitted results are read by a transaction

# Some new terms

**Lost update:** the result of a transaction is overwritten by another transaction

**Dirty read:** uncommitted results are read by a transaction

**Non-repeatable read:** two reads in the same transaction return different results

# Some new terms

**Lost update:** the result of a transaction is overwritten by another transaction

**Dirty read:** uncommitted results are read by a transaction

**Non-repeatable read:** two reads in the same transaction return different results

**Phantom read:** later reads in the same transaction return extra rows


# Serial schedule – No problem

T1: R(A), W(A), R(B), W(B), Abort

T2: R(A), W(A), Commit




# Quiz: Which concurrency problem is this?



Lost update   Dirty read   Non-repeatable read   Phantom read   ??



# Quiz: Which concurrency problem is this?



Lost update   Dirty read   Non-repeatable read   Phantom read   ??


## Quiz: Which concurrency problem is this?

A horizontal timeline with an arrow pointing to the right, labeled "time". Two horizontal lines represent the execution of transactions T1 and T2. T1's operations are: R(A), W(A), and W(B), Commit. T2's operations are: R(A), W(A), W(B), and Commit. The operations are ordered chronologically as they appear on the timeline.

Transaction	Operation	Order
T1	R(A)	1
T2	R(A)	2
T1	W(A)	3
T2	W(A)	4
T1	W(B)	5
T2	W(B)	6
T1	Commit	7
T2	Commit	8

Lost update   Dirty read   Non-repeatable read   Phantom read   ??

# Quiz: Which concurrency problem is this?



Lost update   Dirty read   Non-repeatable read   Phantom read   ??

**Q: How to ensure *correctness* when running concurrent transactions?**

# What does correctness mean?

Transactions should have property of *isolation*, i.e., all operations in a transaction appear to happen together at the same time

# What does correctness mean?

Transactions should have property of *isolation*, i.e., all operations in a transaction appear to happen together at the same time

We need serializability

# Fixing concurrency problems

**Strawman:** Just run transactions serially — prohibitively bad performance

# Fixing concurrency problems

**Strawman:** Just run transactions serially — prohibitively bad performance

**Observation:** Problems only arise when:

1. Two transactions touch the same data
2. At least one of these transactions involves a *write* to the data



# Fixing concurrency problems

**Strawman:** Just run transactions serially — prohibitively bad performance

**Observation:** Problems only arise when:

1. Two transactions touch the same data
2. At least one of these transactions involves a *write* to the data

**Key idea:** Only permit schedules whose effects are guaranteed to be *equivalent* to serial schedules

# Serializability of schedules

Two operations **conflict** if

1. They belong to different transactions
2. They operate on the same data
3. One of them is a **write**

# Serializability of schedules

Two operations **conflict** if

1. They belong to different transactions
2. They operate on the same data
3. One of them is a **write**

Two schedules are **equivalent** if

1. They involve the same transactions and operations
2. All *conflicting* operations are ordered the same way

# Serializability of schedules

Two operations **conflict** if

1. They belong to different transactions
2. They operate on the same data
3. One of them is a **write**

Two schedules are **equivalent** if

1. They involve the same transactions and operations
2. All *conflicting* operations are ordered the same way


A schedule is **serializable** if it is equivalent to a serial schedule

# Testing for serializability

**Intuition:** Swap non-conflicting operations until you reach a serial schedule


# Testing for serializability

**Intuition:** Swap non-conflicting operations until you reach a serial schedule



# Testing for serializability

**Intuition:** Swap non-conflicting operations until you reach a serial schedule



# Testing for serializability

**Intuition:** Swap non-conflicting operations until you reach a serial schedule

T1: R(A), W(A), Commit

T2: R(A), R(B), W(B), Commit

time




# Testing for serializability

**Intuition:** Swap non-conflicting operations until you reach a serial schedule

T1: R(A), W(A), Commit

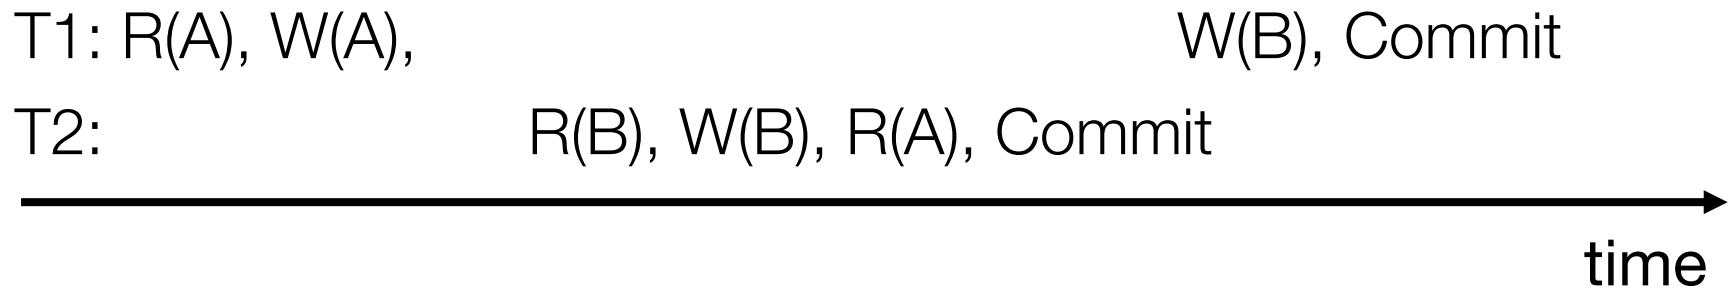
T2: R(A), R(B), W(B) Commit



**Serializable**

# Testing for serializability

**Intuition:** Swap non-conflicting operations until you reach a serial schedule



# Testing for serializability

**Intuition:** Swap non-conflicting operations until you reach a serial schedule


T1: R(A), W(A), W(B), Commit

T2: R(B), W(B), R(A), Commit

time


# Testing for serializability

**Intuition:** Swap non-conflicting operations until you reach a serial schedule



# Testing for serializability

**Intuition:** Swap non-conflicting operations until you reach a serial schedule



**NOT serializable**

# Testing for serializability


Another way to test serializability

- Draw arrows between conflicting operations
- Arrow points in the direction of time
- If no cycles between transactions, the schedule is serializable

# Testing for serializability

Another way to test serializability

- Draw arrows between conflicting operations
- Arrow points in the direction of time
- If no cycles between transactions, the schedule is serializable



# Testing for serializability

Another way to test serializability

- Draw arrows between conflicting operations
- Arrow points in the direction of time
- If no cycles between transactions, the schedule is serializable

T1: R(A), W(A), Commit

T2: R(A), R(B), W(B), Commit






# Testing for serializability

Another way to test serializability

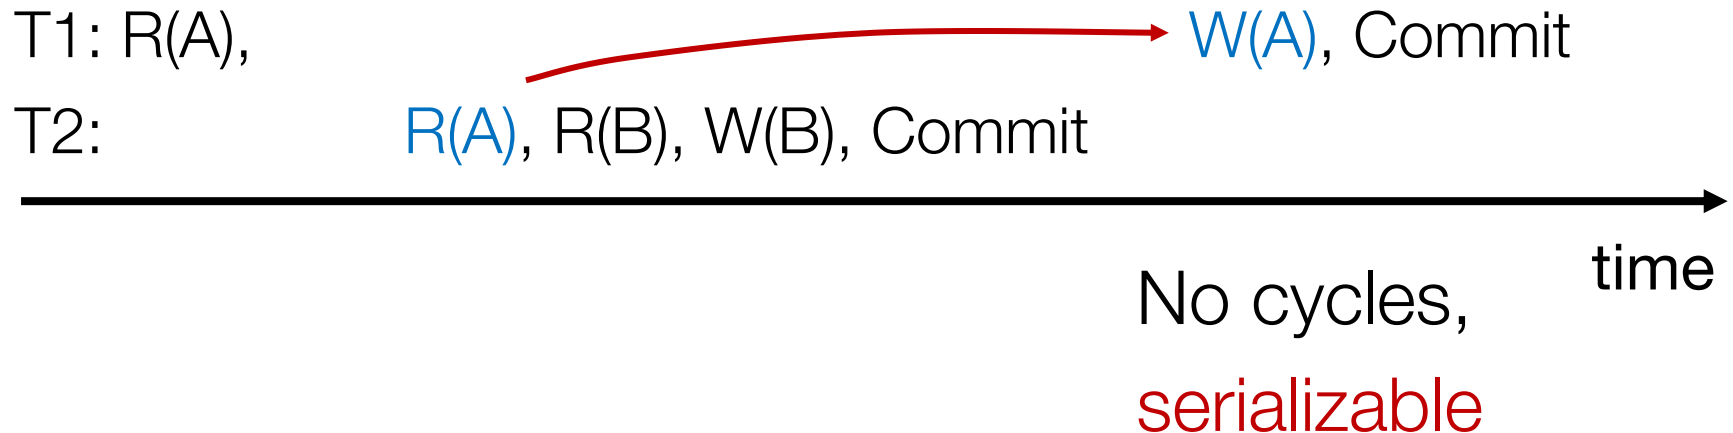
- Draw arrows between conflicting operations
- Arrow points in the direction of time
- If no cycles between transactions, the schedule is serializable



# Testing for serializability

Another way to test serializability

- Draw arrows between conflicting operations
- Arrow points in the direction of time
- If no cycles between transactions, the schedule is serializable



# Testing for serializability

Another way to test serializability

- Draw arrows between conflicting operations
- Arrow points in the direction of time
- If no cycles between transactions, the schedule is serializable

T1: R(A), W(A), W(B), Commit


T2: R(B), W(B), R(A), Commit



# Testing for serializability

Another way to test serializability


- Draw arrows between conflicting operations
- Arrow points in the direction of time
- If no cycles between transactions, the schedule is serializable



# Testing for serializability

Another way to test serializability

- Draw arrows between conflicting operations
- Arrow points in the direction of time
- If no cycles between transactions, the schedule is serializable



# Testing for serializability

Another way to test serializability

- Draw arrows between conflicting operations
- Arrow points in the direction of time
- If no cycles between transactions, the schedule is serializable



# Linearizability vs. Serializability

- **Linearizability**: a guarantee about **single** operations on **single** objects
  - Once write completes, all later reads (by wall clock) should reflect that write
- **Serializability** is a guarantee about **transactions** over one or more objects
  - Doesn't impose real-time constraints
- Linearizability + serializability = **strict serializability**
  - Transaction behavior equivalent to some serial execution
    - And that serial execution agrees with real-time

# Lock-based concurrency control


- **Big Global Lock:** Results in a **serial** transaction schedule at the **cost of performance**
- **2PL: Two-phase locking with finer-grain locks:**
  - **Growing phase** when txn acquires locks
  - **Shrinking phase** when txn releases locks (typically commit)
  - Allows txns to execute concurrently, improving performance



# 2PL


- 2PL guarantees **serializability** by disallowing cycles between txns
- There could be dependencies in the waits-for graph among txns waiting for locks:
  - Edge from T2 to T1 means T1 acquired lock first and T2 has to wait
  - Edge from T1 to T2 means T1 acquired lock first and T1 has to wait
  - Cycles mean **DEADLOCK**, and in that case 2PL won't proceed

# 2PL




Deal with deadlocks by aborting one of the two txns (e.g., detect with timeout)

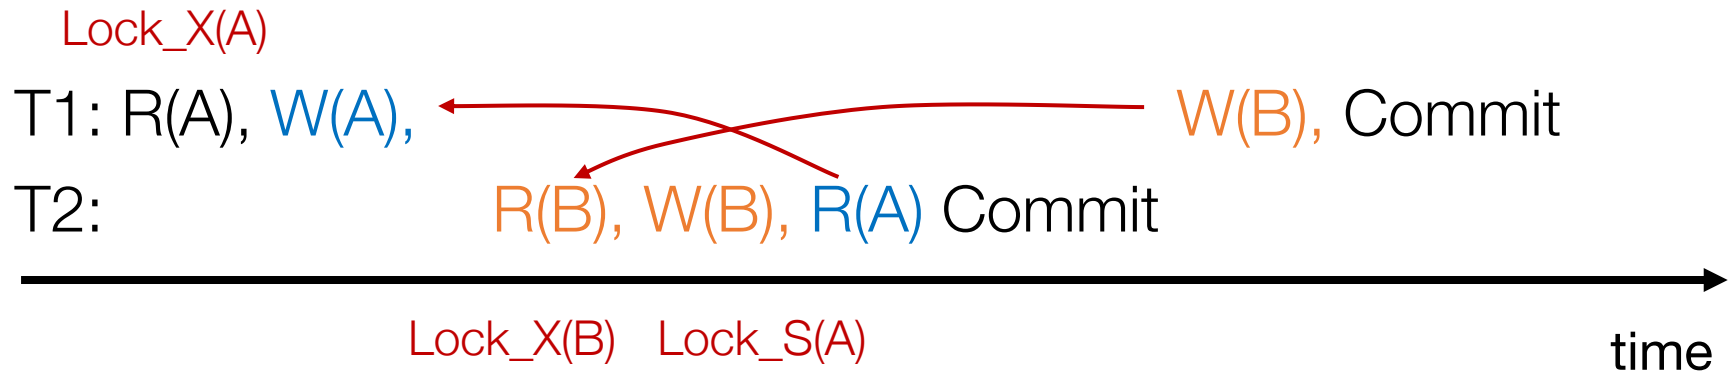
# 2PL



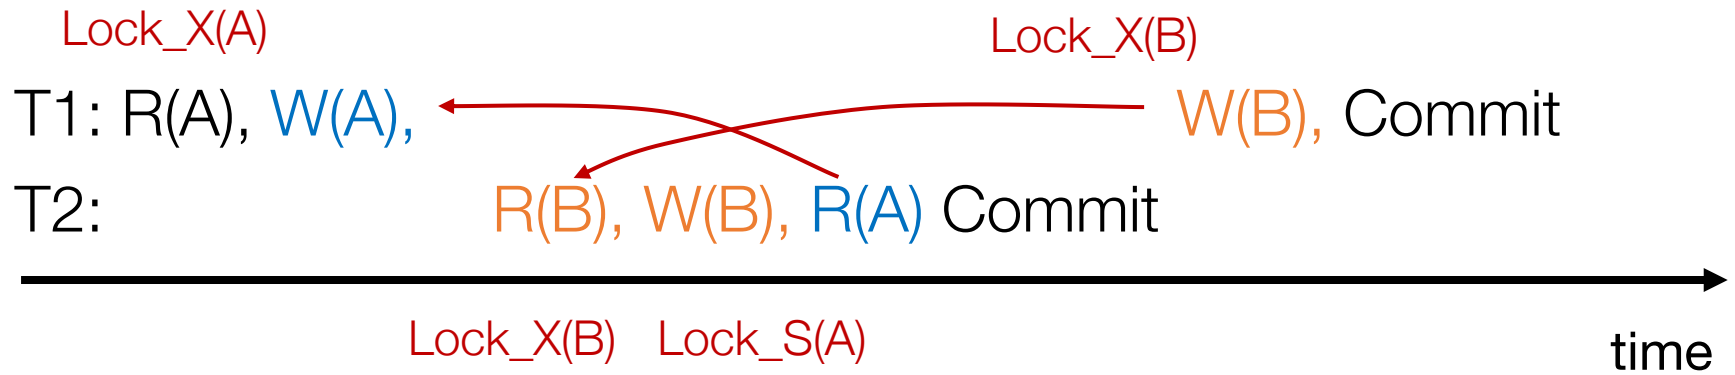
# 2PL




# 2PL




# 2PL



# 2PL



# 2PL



Deal with deadlocks by aborting one of the two txns (e.g., detect with timeout)



# 2PL: Releasing locks too soon?

What if we release the lock as soon as we can?

# 2PL: Releasing locks too soon?

What if we release the lock as soon as we can?



# 2PL: Releasing locks too soon?

What if we release the lock as soon as we can?




# 2PL: Releasing locks too soon?

What if we release the lock as soon as we can?




# 2PL: Releasing locks too soon?

What if we release the lock as soon as we can?




# 2PL: Releasing locks too soon?

What if we release the lock as soon as we can?



# 2PL: Releasing locks too soon?


What if we release the lock as soon as we can?



Rollback of T1 requires rollback of T2, since T2 reads a value written by T1

# 2PL: Releasing locks too soon?

What if we release the lock as soon as we can?



Rollback of T1 requires rollback of T2, since T2 reads a value written by T1

**Cascading aborts:** the rollback of one txn causes rollback of another



# Strict 2PL

- Release locks at the end of the transaction
- Variant of 2PL implemented by most DBs in practice

**Q: What if access patterns rarely,  
if ever, conflict?**

# Today

- Optimistic concurrency control (OCC)
  - Be optimistic, or opportunistic that conflicts rarely happen

# Be optimistic!

- Goal: Low overhead for non-conflicting txns
- Assume success!
  - Process transaction as if would succeed
  - Check for serializability only at commit time
  - If fails, abort transaction
- Optimistic Concurrency Control (OCC)
  - **Higher performance** when few conflicts vs. locking
  - **Lower performance** when many conflicts vs. locking

# OCC: Three-phase approach

- **Begin:** Record timestamp marking the transaction's beginning

# OCC: Three-phase approach

- **Begin:** Record timestamp marking the transaction's beginning
- **Modify phase:**
  - Txn can read values of committed data items
  - Updates only to local copies (versions) of items (in DB cache)

# OCC: Three-phase approach

- **Begin:** Record timestamp marking the transaction's beginning
- **Modify phase:**
  - Txn can read values of committed data items
  - Updates only to local copies (versions) of items (in DB cache)
- **Validate** phase

# OCC: Three-phase approach

- **Begin:** Record timestamp marking the transaction's beginning
- **Modify phase:**
  - Txn can read values of committed data items
  - Updates only to local copies (versions) of items (in DB cache)
- **Validate phase**
- **Commit phase**
  - If validates, transaction's updates applied to DB
  - Otherwise, transaction restarted
  - Care must be taken to avoid "TOCTTOU" issues



# OCC: Three-phase approach

- **Begin:** Record timestamp marking the transaction's beginning

- **Modify phase:**

Execute optimistically!

- Txn can read values of committed data items
- Updates only to local copies (versions) of items (in DB cache)

- **Validate phase**

- **Commit phase**

- If validates, transaction's updates applied to DB
- Otherwise, transaction restarted
- Care must be taken to avoid "TOCTTOU" issues

# OCC: Three-phase approach

- **Begin:** Record timestamp marking the transaction's beginning

- **Modify phase:**

Execute optimistically!

- Txn can read values of committed data items
- Updates only to local copies (versions) of items (in DB cache)

- **Validate phase**


- **Commit phase**

These should happen together!

- If validates, transaction's updates applied to DB
- Otherwise, transaction restarted
- Care must be taken to avoid "TOCTTOU" issues

# OCC: Why validation is necessary!

- New txn creates shadow copies of P and Q
- P and Q's copies at inconsistent state



When **commits** txn updates, create new versions at some timestamp  $t$

# OCC: Validate phase

- Transaction is about to commit. System must ensure:
  - **Initial consistency:** Versions of accessed objects at start consistent
  - **No conflicting concurrency:** No other txn has committed an operation at object that conflicts with one of this txn's invocations
- Consider transaction 1: For all other txns N either committed or in validation phase, one of the following holds:
  - A. N completes commit before 1 starts modify
  - B. 1 starts commit after N completes commit, and ReadSet 1 and WriteSet N are disjoint
  - C. Both ReadSet 1 and WriteSet 1 are disjoint from WriteSet N, and N completes modify phase
- When validating 1, first check (A), then (B), then (C). If all fail, validation fails and 1 aborted

# Atomic commit for OCC

- Use **two-phase commit (2PC)** to achieve atomic commit (validate + commit writes)
- Recall 2PC protocol:
  1. Coordinator sends *prepare* messages to all nodes, other nodes vote *yes* or *no*
    - a. If all nodes accept, proceed
    - b. If any node declines, abort
  2. Coordinator sends *commit* or *abort* messages to all nodes, and all nodes act accordingly

# Atomic commit for OCC

- **Execute optimistically:** Read committed values, write changes locally

• **Validate:** Check if data has changed since original read | Phase 1

• **Commit (Write):** Commit if no change, else abort | Phase 2

# Atomic commit for OCC

- **Execute optimistically:** Read committed values, write changes locally

• **Validate:** Check if data has changed since original read | Phase 1

• **Commit (Write):** Commit if no change, else abort | Phase 2

- **Phase 1:** send *prepare* to each shard: include buffered write + original reads for that shard
  - Shards **validate reads and acquire locks** (exclusive for write locations, shared for read locations)
  - If this succeeds, respond with *yes*; else respond with *no*

# Atomic commit for OCC

- **Execute optimistically:** Read committed values, write changes locally

• **Validate:** Check if data has changed since original read | Phase 1

• **Commit (Write):** Commit if no change, else abort | Phase 2

- **Phase 1:** send *prepare* to each shard: include buffered write + original reads for that shard
  - Shards **validate reads and acquire locks** (exclusive for write locations, shared for read locations)
  - If this succeeds, respond with *yes*; else respond with *no*
- **Phase 2:** collect votes, send result (*abort* or *commit*) to all shards
  - If commit, **shards apply buffered writes**
  - All shards release locks



# Two ways of implementing serializability: 2PL, OCC

- 2PL (**pessimistic**):
  - Assume conflict, always lock
  - High overhead for non-conflicting txn
  - Must check for deadlock
- OCC (**optimistic**):
  - Assume no conflict
  - Low overhead for low-conflict workloads (but high for high-conflict workloads)
  - Ensure correctness by aborting txns if conflict occurs

Lock_X(A) <granted>	
Read(A)	Lock_S(A)
A := A-50	↓
Write(A)	↓
Unlock(A)	<granted>
	Read(A)
	Unlock(A)
	Lock_S(B) <granted>
Lock_X(B)	
↓	Read(B)
<granted>	Unlock(B)
Read(B)	
B := B +50	
Write(B)	
Unlock(B)	

Is this a 2PL schedule?

No

Is this a serializable schedule?

No

<b>Lock_X(A) &lt;granted&gt;</b>	
<b>Read(A)</b>	<b>Lock_S(A)</b>
<b>A := A-50</b>	
<b>Write(A)</b>	
<b>Lock_X(B) &lt;granted&gt;</b>	
<b>Unlock(A)</b>	<b>&lt;granted&gt;</b>
	<b>Read(A)</b>
	<b>Lock_S(B)</b>
<b>Read(B)</b>	
<b>B := B +50</b>	
<b>Write(B)</b>	
<b>Unlock(B)</b>	<b>&lt;granted&gt;</b>
	<b>Unlock(A)</b>
	<b>Read(B)</b>
	<b>Unlock(B)</b>

Is this a 2PL schedule?

Yes, and it is serializable

Is this a Strict 2PL schedule?

No, cascading aborts possible

Lock_X(A) <granted>	
Read(A)	Lock_S(A)
A := A-50	
Write(A)	
Lock_X(B) <granted>	
Read(B)	
B := B +50	
Write(B)	
Unlock(A)	
Unlock(B)	<granted>
	Read(A)
	Lock_S(B) <granted>
	Read(B)
	Unlock(A)
	Unlock(B)

Is this a 2PL schedule?

Yes, and it is serializable

Is this a Strict 2PL schedule?

Yes, cascading aborts not possible