

Resilient Distributed Datasets: Spark

CS 475: Concurrent & Distributed Systems (Fall 2021)

Lecture 16

Yue Cheng

Some material taken/derived from:

- Matei Zaharia's NSDI'12 talk slides.
- Utah CS6450 by Ryan Stutsman.

Licensed for use under a Creative Commons Attribution-NonCommercial-ShareAlike 3.0 Unported License.

What's good with MapReduce

- Scaled analytics to thousands of machines
- Eliminated fault-tolerance as a concern

Problems with MapReduce

- Scaled analytics to thousands of machines
- Eliminated fault-tolerance as a concern
- **Not very expressive**
 - Iterative algorithms
(PageRank, Logistic Regression, Transitive Closure)
 - Interactive and ad-hoc queries
(Interactive Log Debugging)
- Lots of specialized frameworks
 - Pregel, GraphLab, PowerGraph, DryadLINQ, HaLoop...



Sharing data between iterations/ops

- Only way to share data between iterations / phases is through shared storage
 - **Slow!**
- Allow operations to feed data to one another
 - Ideally, through memory instead of disk-based storage



Sharing data between iterations/ops

- Only way to share data between iterations / phases is through shared storage
 - **Slow!**
- Allow operations to feed data to one another
 - Ideally, through memory instead of disk-based storage
- Need the “chain” of operations to be exposed to make this work
- Also, does this break the MR fault-tolerance scheme?
 - Retry and Map or Reduce task since idempotent

Examples





Examples





Slow due to replication and disk I/O,
but necessary for fault tolerance

Goal: In-memory data sharing



Goal: In-memory data sharing



10-100× faster than network/disk, but how to get FT?

Challenges

- How to design a distributed memory abstraction that is both **fault-tolerant** and **efficient**?

Challenges

- How to design a distributed memory abstraction that is both **fault-tolerant** and **efficient**?
- Existing storage systems allow **fine-grained** mutation to state
 - In-memory key-value stores
 - Requires replicating data or logs across nodes for fault tolerance
 - Costly for data-intensive apps
 - 10-100x slower than memory write
 - They also require costly on-the-fly replication for mutations

Challenges

- How to design a distributed memory abstraction that is both **fault-tolerant** and **efficient**?
- Existing storage systems allow **fine-grained** mutation to state

Insight: leverage similar coarse-grained approach that transforms whole dataset per operation, like MapReduce (batch processing)

- 10-100x slower than memory write
- They also require costly on-the-fly replication for mutations

Solution: Resilient Distributed Datasets (RDDs)

- Restricted form of distributed shared memory
 - **Immutable**, partitioned collections of records
 - Can only be built through **coarse-grained**, deterministic **transformations** (map, filter, join, ...)
- Efficient fault recovery using **lineage**
 - Log **one operation** to apply to many elements
 - Recompute lost partitions on failure
 - No cost if nothing fails

Spark programming interface

- Scala API, exposed within interpreter as well
- RDDs
- Transformations on RDDs ($RDD_1 \rightarrow RDD_2$)
- Actions on RDDs ($RDD \rightarrow \text{output}$)
- Control over RDD partitioning (how items are split over nodes)
- Control over RDD persistence (in memory, on disk, or recompute on loss)

Transformations

Transformations (define a new RDD)	map filter sample groupByKey reduceByKey sortByKey	flatMap union join cogroup cross mapValues
---------------------------------------	---	---

RDDs in terms of Scala types → Scala semantics at workers

Transformations are **lazy “thunks”**; cause no cluster action

Actions

Actions (return a result to driver program)	collect reduce count save lookupKey
--	---

Consumes an RDD to **produce** output
either to storage (save), or
to interpreter/Scala (count, collect, reduce)

Causes RDD lineage chain to get executed on the cluster to
produce the output
(for any missing pieces of the computation)


Interactive debugging

```
lines = textFile("hdfs://foo.log")
errors = lines.filter(
    _.startsWith("ERROR")
errors.persist()
```

Interactive debugging

```
lines = textFile("hdfs://foo.log")
errors = lines.filter(
    _.startsWith("ERROR")
)
errors.persist()
```


```
errors.count()
```



Interactive debugging

```
lines = textFile("hdfs://foo.log")
errors = lines.filter(
    _.startsWith("ERROR")
)
errors.persist()

errors.count()
```




Interactive debugging

```
lines = textFile("hdfs://foo.log")
errors = lines.filter(
    _.startsWith("ERROR")
errors.persist()
```

```
errors.count()
```

```
errors.filter(
    _.contains("MySQL"))
```




Interactive debugging

```
lines = textFile("hdfs://foo.log")
errors = lines.filter(
    _.startsWith("ERROR")
errors.persist()
```

```
errors.count()
```

```
errors.filter(
    _.contains("MySQL")).count()
```




Interactive debugging

```
lines = textFile("hdfs://foo.log")
errors = lines.filter(
    _.startsWith("ERROR")
errors.persist()

errors.count()

errors.filter(
    _.contains("MySQL")).count()
```




Interactive debugging

```
lines = textFile("hdfs://foo.log")
errors = lines.filter(
    _.startsWith("ERROR"))
errors.persist()
```

```
errors.count()
```

```
errors.filter(
    _.contains("MySQL")).count()
```




Interactive debugging

```
lines = textFile("hdfs://foo.log")
errors = lines.filter(
    _.startsWith("ERROR")
errors.persist()
```

```
errors.count()
```

```
errors.filter(
    _.contains("MySQL")).count()
errors.filter(
    _.contains("HDFS"))
```




Interactive debugging

```
lines = textFile("hdfs://foo.log")
errors = lines.filter(
    _.startsWith("ERROR"))
errors.persist()
```

```
errors.count()
```

```
errors.filter(
    _.contains("MySQL")).count()
errors.filter(
    _.contains("HDFS"))
    _.map(_.split("\t")(3))
```




Interactive debugging

```
lines = textFile("hdfs://foo.log")
errors = lines.filter(
    _.startsWith("ERROR"))
errors.persist()
```

```
errors.count()
```

```
errors.filter(
    _.contains("MySQL")).count()
errors.filter(
    _.contains("HDFS"))
    _.map(_.split("\t")(3))
    .collect()
```




Interactive debugging

```
lines = textFile("hdfs://foo.log")
errors = lines.filter(
    _.startsWith("ERROR"))
errors.persist()
```

```
errors.count()
```

```
errors.filter(
    _.contains("MySQL")).count()
errors.filter(
    _.contains("HDFS"))
    _.map(_.split("\t")(3))
    .collect()
```




Interactive debugging

```
lines = textFile("hdfs://foo.log")
errors = lines.filter(
    _.startsWith("ERROR"))
errors.persist()
```

```
errors.count()
```

```
errors.filter(
    _.contains("MySQL")).count()
errors.filter(
    _.contains("HDFS"))
    _.map(_.split("\t")(3))
    .collect()
```




Interactive debugging

```
lines = textFile("hdfs://foo.log")
errors = lines.filter(
    _.startsWith("ERROR"))
errors.persist()
```

```
errors.count()
```

```
errors.filter(
    _.contains("MySQL")).count()
errors.filter(
    _.contains("HDFS"))
    _.map(_.split("\t")(3))
    .collect()
```




`persist()`

- Not an action and not a transformation
- A scheduler hint
- Tells which RDDs the Spark scheduler should materialize and whether in memory or storage
- Gives the user control over reuse/recompute/recovery tradeoffs

persist()

- Not an action and not a transformation
- A scheduler hint
- Tells which RDDs the Spark scheduler should materialize and whether in memory or storage
- Gives the user control over reuse/recompute/recovery tradeoffs
- **Q:** If persist() asks for the materialization of an RDD why isn't it an action?



Lineage graph of RDDs




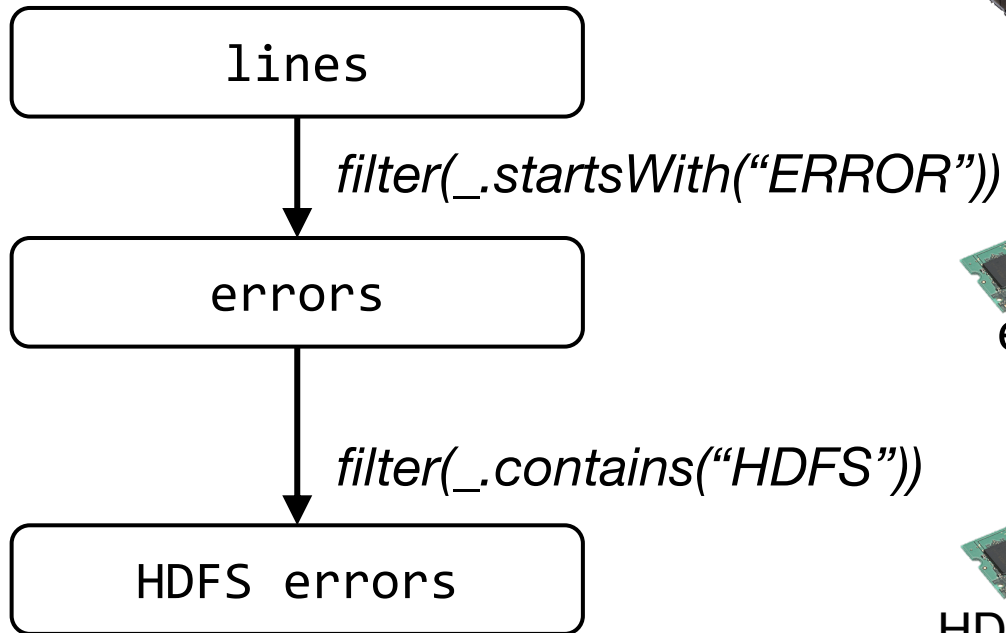
lines



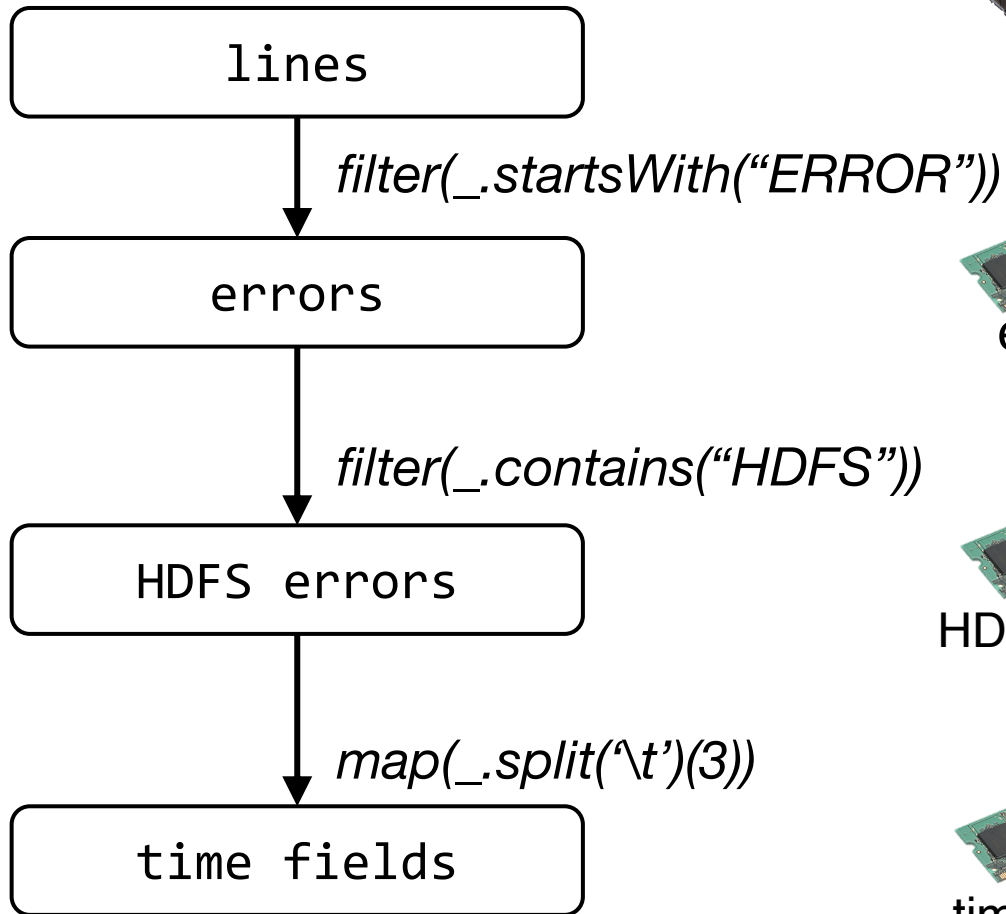
Lineage graph of RDDs





Lineage graph of RDDs




Lineage graph of RDDs




lines



errors



HDFS errors




time fields




Narrow & wide dependencies

Narrow Dependencies:



Wide Dependencies:



Narrow: each parent partition used by at most one child partition
(can partition on one machine)

Wide: multiple child partitions depend on one parent partition

Must stall for all parent data, loss of child requires whole parent RDD (not just a small # of partitions)


Task scheduler

Dryad-like DAGs

Pipelines functions within a stage

Locality & data reuse aware

Partitioning-aware to avoid shuffles



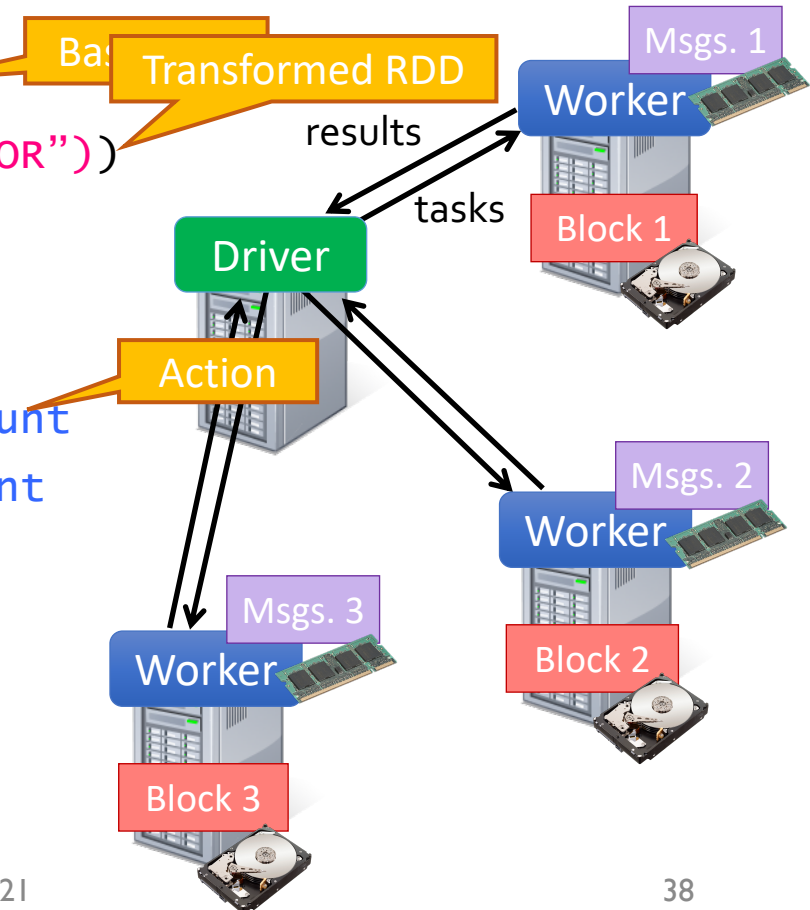
Interactive debugging (control and data flow)

Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://...")
errors = lines.filter(_.startsWith("ERROR"))
messages = errors.map(_.split('\t')(2))
messages.persist()

messages.filter(_.contains("MySQL")).count
messages.filter(_.contains("HDFS")).count
```

Result: scaled to 1 TB data in 5-7 sec
(vs 170 sec for on-disk data)

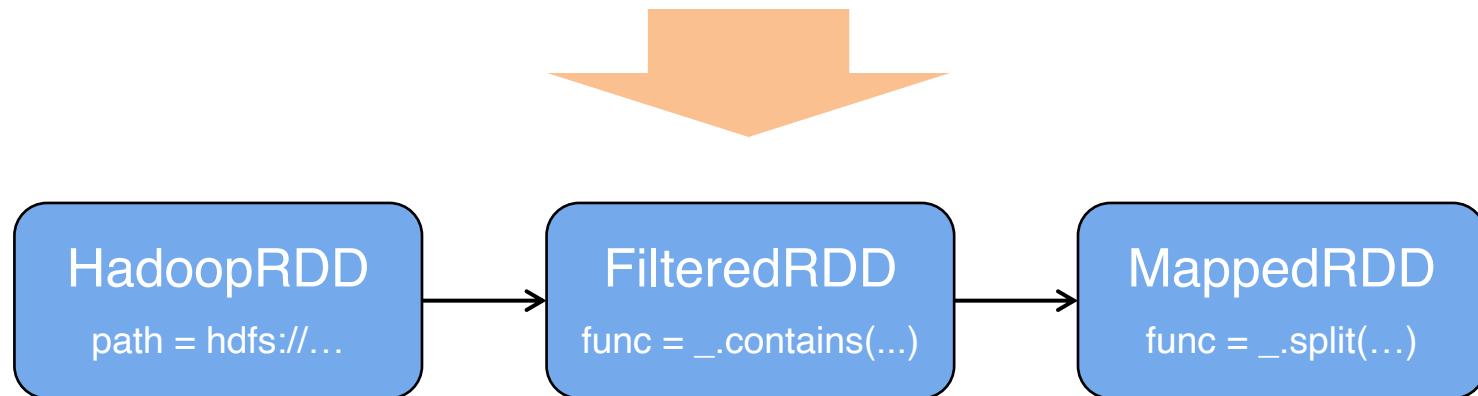


Fault recovery

- RDDs track the graph of transformations that built them (their *lineage*) to rebuild lost data

E.g.:

```
messages = textFile(...).filter(_.contains("error"))  
                        .map(_.split('\t')(2))
```




Fault recovery

- RDDs track the graph of transformations that built them (their *lineage*) to rebuild lost data

E.g.:

```
messages = textFile(...).filter(_.contains("error"))  
                        .map(_.split('\t')(2))
```




Fault recovery


- RDDs track the graph of transformations that built them (their *lineage*) to rebuild lost data

E.g.:

```
messages = textFile(...).filter(_.contains("error"))  
                        .map(_.split('\t')(2))
```



Fault recovery results



Example: PageRank

1. Start each page with a rank of 1
2. On each iteration, update each page's rank to

$$\sum_{i \in \text{neighbors}} \text{rank}_i / |\text{neighbors}_i|$$

links = // RDD of (url, neighbors) pairs

ranks = // RDD of (url, rank) pairs

```
for (i <- 1 to ITERATIONS) {  
  ranks = links.join(ranks).flatMap {  
    (url, (links, rank)) =>  
      links.map(dest => (dest, rank/links.size))  
  }.reduceByKey(_ + _)  
}
```

Example: PageRank

1. Start each page with a rank of 1
2. On each iteration, update each page's rank to

$$\sum_{i \in \text{neighbors}} \text{rank}_i / |\text{neighbors}_i|$$

RDD[(URL, Seq[URL])]

links = // RDD of (url, neighbors) pairs

ranks = // RDD of (url, rank) pairs ← RDD[(URL, Rank)]

for (i <- 1 to ITERATIONS) { ← RDD[(URL, (Seq[URL], Rank))]

 ranks = links.join(ranks).flatMap {

 (url, (links, rank)) =>

 links.map(dest => (dest, rank/links.size))

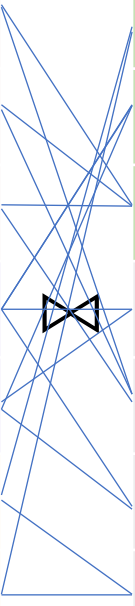
 }.reduceByKey(_ + _)

} ← For each neighbor in links emits (URL, RankContrib)

Reduce to RDD[(URL, Rank)]


Join (\bowtie)

Alice	5	\bowtie	Alice	F	=	Alice	5	F
Bob	6		Bob	M		Bob	6	M
Claire	4		Claire	F		Claire	4	F

A	5		C	5
A	2		B	2
A	3		A	3
B	4		B	4
B	1		A	1
C	6		B	6
C	8		C	8


If partitioning doesn't match, then need to reshuffle to match pairs. Same problem in `reduce()` for MapReduce.

Optimizing placement



- Links & ranks repeatedly joined
- Can *co-partition* them (e.g. hash both on URL) to avoid shuffles
- Can also use app knowledge, e.g., hash on DNS name
- `links = links.partitionBy(new URLPartitioner())`

Optimizing placement



- Links & ranks repeatedly joined
- Can *co-partition* them (e.g. hash both on URL) to avoid shuffles
- Can also use app knowledge, e.g., hash on DNS name
- `links = links.partitionBy(new URLPartitioner())`

Q: Where might we have placed **`persist()`**?


Co-partitioning example

Co-partitioning can avoid shuffle on join

But, fundamentally a shuffle on **reduceByKey**

Optimization: custom partitioner on domain

PageRank performance



* Figure 10a: 30 machines on 54 GB of Wikipedia data computing PageRank

Tradeoff space

